

# **Software Security**

Prof. Leonardo B. Oliveira

# Agenda

- C language
- Memory Management
- Attacks
- Countermeasures
- Final Remarks

# Look at this code

```
#include <stdio.h>

int main() {
    int array[5] = {1, 2, 3, 4, 5};
    printf("%d\n", array[5]);
    return 0;
}
```

- What is this function doing?

# Look at this code

```
#include <stdio.h>

int main() {
    int array[5] = {1, 2, 3, 4, 5};
    printf("%d\n", array[5]);
    return 0;
}
```

- What is this function doing?
- What's going to happen if we compile it?

# Look at this code

```
#include <stdio.h>

int main() {
    int array[5] = {1, 2, 3, 4, 5};
    printf("%d\n", array[5]);
    return 0;
}
```

- What is this function doing?
- What's going to happen if we compile it?
- What's going to happen if we run it?

# Demo

Access past the end of an array

# Look at this code

```
#include <stdio.h>

int main() {
    int array[5] = {1, 2, 3, 4, 5};
    printf("%d\n", array[5]);
    return 0;
}
```

```
security_course$ cc test.cpp
security_course$ ./a.out
32767
security_course$ █
```

# Look at this code

```
#include <stdio.h>

int main() {
    int array[5] = {1, 2, 3, 4, 5};
    printf("%d\n", array[5]);
    return 0;
}
```

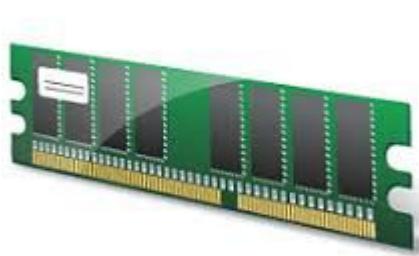
- In C there's no array bound-checking
- An illegal array access can be performed
- Like accessing the 6th element of a 5-element array
- One can read (or write) past the end of the array!
- **And then bad things might happen**

# Agenda

- C language
- **Memory Management**
- Attacks
- Countermeasures
- Final Remarks

# Memory Management: Concepts

- If you know how something works, you can protect yourself against it
  - This holds true for computer stuff as well!
- You must be familiar with computer architecture, object files, and the runtime stack



# Memory Organization

For each program, the OS allocates 3 different memory segments

# Memory Organization

For each program, the OS allocates 3 different memory segments

Q: What are they?

# Memory Organization

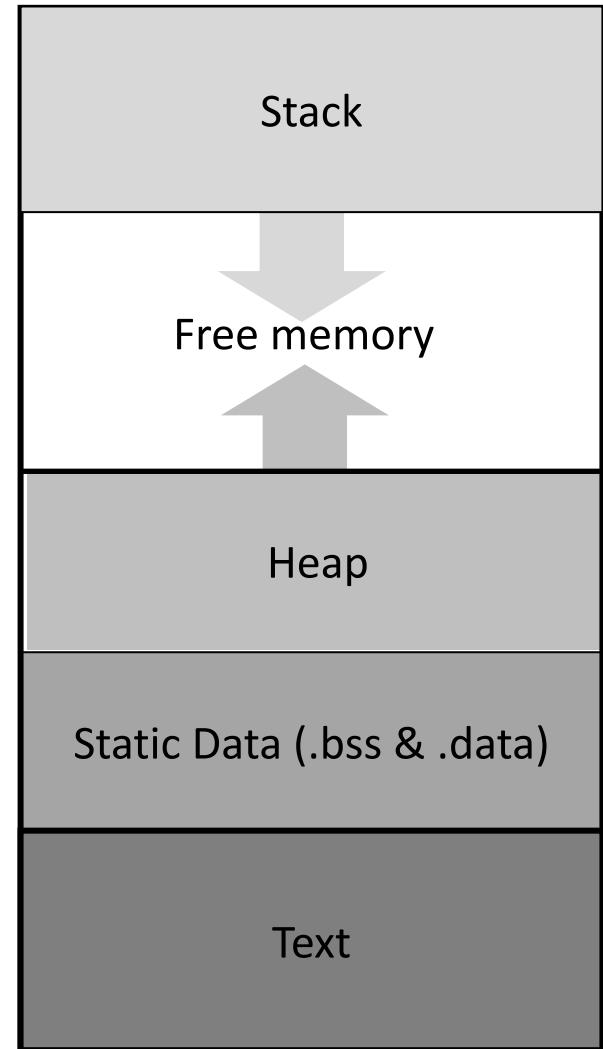
Stack: stores function's local variables and its own control information

For each program, the OS allocates 3 different memory segments

Q: What are they?

Data segment: stores both dynamic (heap) and also static (.bss and .data) data

Text segment: program instructions in the form of machine code



# Byte Ordering

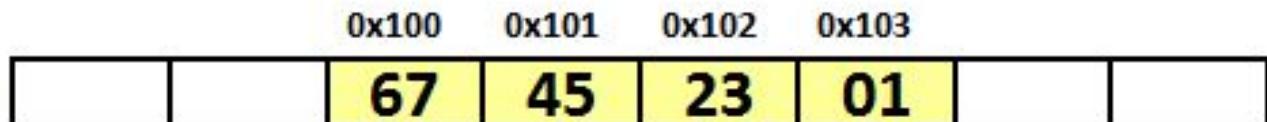
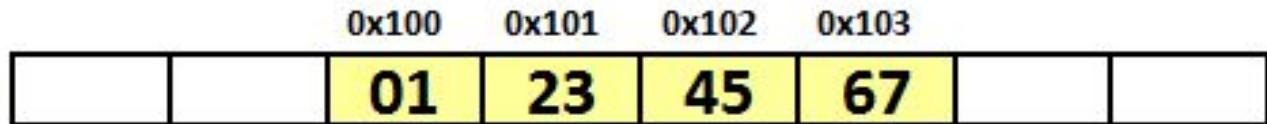
- Big endian – most significant (big) byte in smallest address
  - E.g. Sun
- Little endian – least significant (little) byte in smallest address
  - E.g. IA32

# Byte Ordering

- Big endian – most significant (big) byte in smallest address
  - E.g. Sun
- Little endian – least significant (little) byte in smallest address
  - E.g. IA32
- How would the word 0x01234567 be placed at address 0x100?

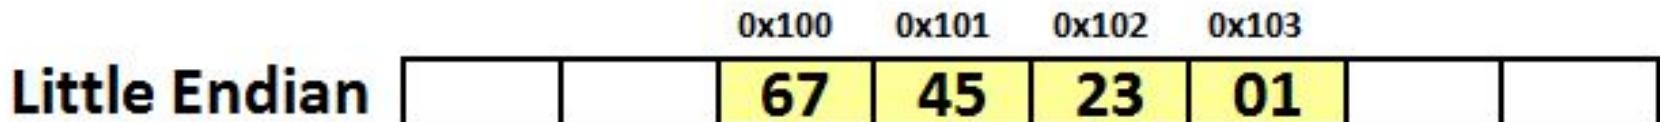
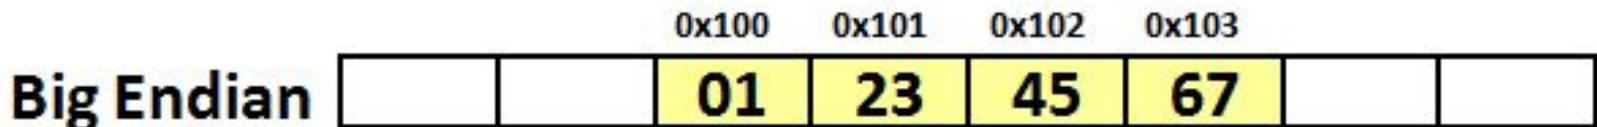
# Byte Ordering

- Big endian – most significant (big) byte in smallest address
  - E.g. Sun
- Little endian – least significant (little) byte in smallest address
  - E.g. IA32
- Q: How would the word 0x01234567 be placed at address 0x100?
  - Q: Which is big or little endian?



# Byte Ordering

- Big endian – most significant (big) byte in smallest address
  - E.g. Sun
- Little endian – least significant (little) byte in smallest address
  - E.g. IA32
- Q: How would the word 0x01234567 be placed at address 0x100?
  - Q: Which is big or little endian?



# Stack Alignment

- Easy concept, but fairly important
- Word alignment requires that words begin in multiple of  $n$ , e.g.:
  - $n = 4$  for 32-bit architecture
  - $n = 8$  for 64-bit architecture

32-bit Words	64-bit Words	Bytes	Addr.
Addr = 0000	Addr = 0000		0000
Addr = 0004			0001
Addr = 0008			0002
Addr = 0012	Addr = 0008		0003
			0004
			0005
			0006
			0007
			0008
			0009
			0010
			0011
			0012
			0013
			0014
			0015

# Stack Alignment

```
struct {  
    char x; // 1 byte  
    int y; // 4 bytes  
    char z; // 1 byte  
    int w; // 4 bytes  
} foo;
```

- Q: What's the size of foo?

# Stack Alignment

```
struct {  
    char x; // 1 byte  
    int y; // 4 bytes  
    char z; // 1 byte  
    int w; // 4 bytes  
} foo;
```

- Q: What's the size of foo?
- A: 10 bytes?

# Stack Alignment

```
#include <stdio.h>

struct {
    char x; // 1 byte
    int y; // 4 bytes
    char z; // 1 byte
    int w; // 4 bytes
} foo;

int main() {
    printf("sizeof(foo) == %lu\n", sizeof(foo));
    return 0;
}
```

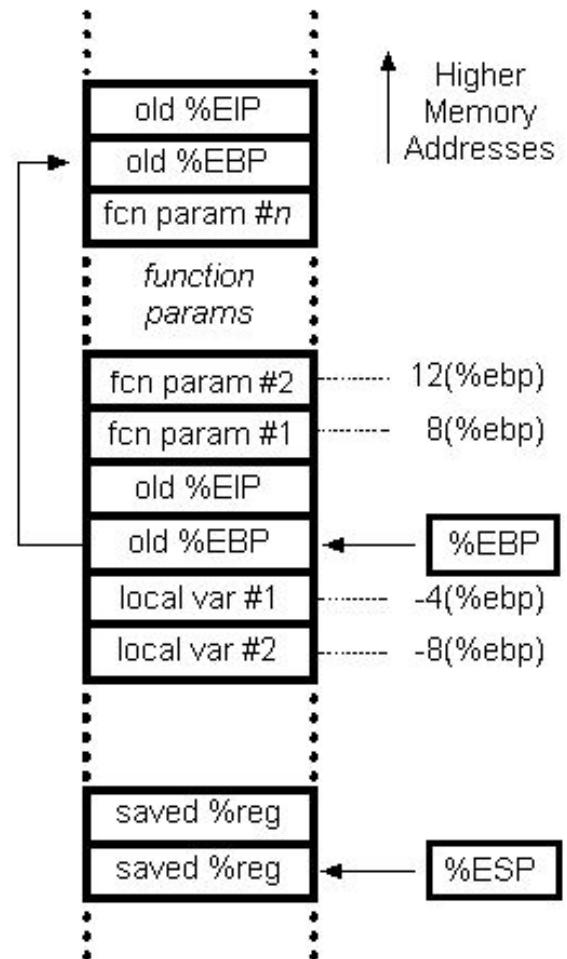
```
security_course$ clang test.c
security_course$ ./a.out
sizeof(foo) == 16
```

# Demo

struct size

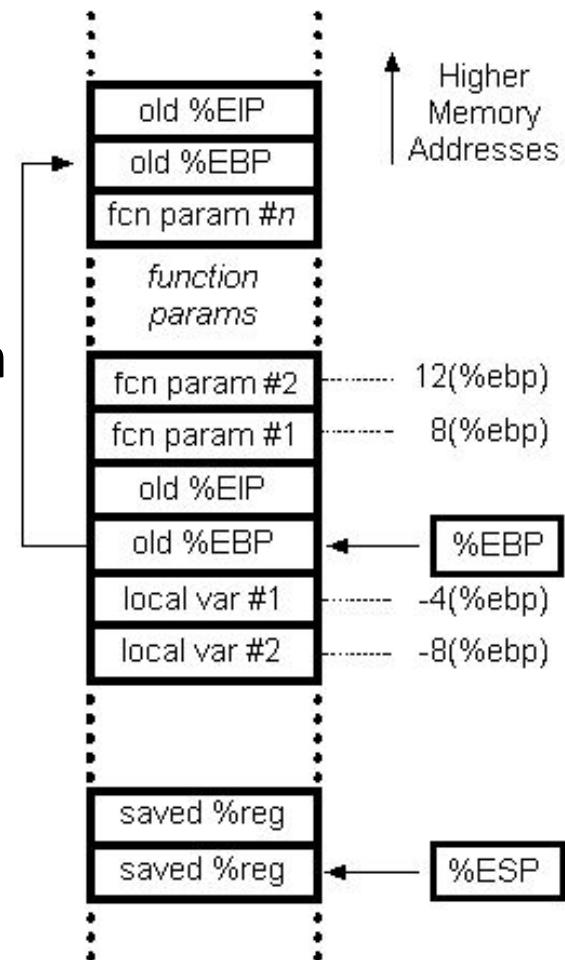
# Functions & The Stack

- Q: What's the purpose of a stack?



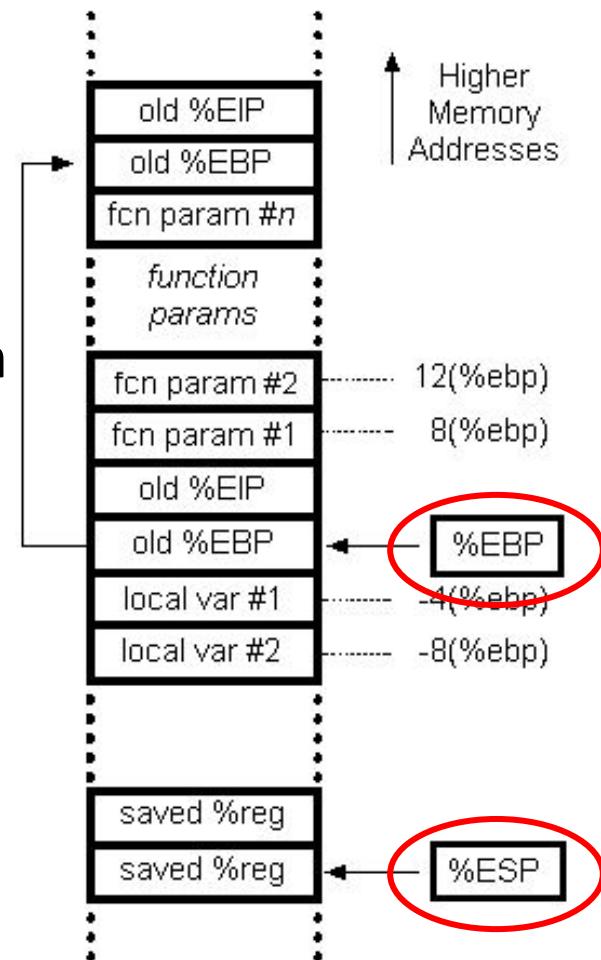
# Functions & The Stack

- Q: What's the purpose of a stack?
- A: To implement the dynamics of a function efficiently
  - Change the execution flow of a program
  - Execute some instructions
  - Return control to the point where the function was called



# Functions & The Stack

- Q: What's the purpose of a stack?
- A: To implement the dynamics of a function efficiently
  - Change the execution flow of a program
  - Execute some instructions
  - Return control to the point where the function was called
- Especial registers
  - ESP
  - EBP



# Special Registers

- ESP
  - ESP points to the top of the stack
  - In the IA32 architecture, it points to the latest used address

# Special Registers

- ESP
  - ESP points to the top of the stack
  - In the IA32 architecture, it points to the latest used address
- EBP
  - AKA, *frame or base pointer*
  - It is often used to calculate and address based on another address, i.e., EBP plus an offset

# Function Dynamics

- Consider the following code

```
#include <stdio.h>

void foo(int x) {
    printf("Hello World: %d\n", x);
}

int main() {
    foo(1);
    return 0;
}
```



# Function Dynamics

- Consider the following code

```
#include <stdio.h>

void foo(int x) {
    printf("Hello World: %d\n", x);
}

int main() {
     foo(1);
    return 0;
}
```

# Function Dynamics

- Consider the following code

```
#include <stdio.h>

void foo(int x) {
     printf("Hello World: %d\n", x);
}

int main() {
    foo(1);
    return 0;
}
```

# Function Dynamics

- Consider the following code

```
#include <stdio.h>

void foo(int x) {
    printf("Hello World: %d\n", x);
}

int main() {
    foo(1);
    return 0;
}
```



If you wanna get software security right you really need to understand this function dynamics!



# Right before a function call

- Q: What happens in between a function call and its actual execution?

# Right before a function call

- Q: What happens in between a function call and its actual execution?
- A: Whenever a function call is made, the following instructions are carried out
  1. Push function's arguments, in reverse order, onto the stack
  2. Push the return address (RET) onto the stack

# Right before a function call

- Q: What happens in between a function call and its actual execution?
  - A: Whenever a function call is made, the following instructions are carried out
    1. Push function's arguments, in reverse order, onto the stack
    2. Push the return address (RET) onto the stack
- Q: What is RET and why are we pushing it?
- A: Yes, so we can come back to where we came from

...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp
```

...

**leave**

**ret**

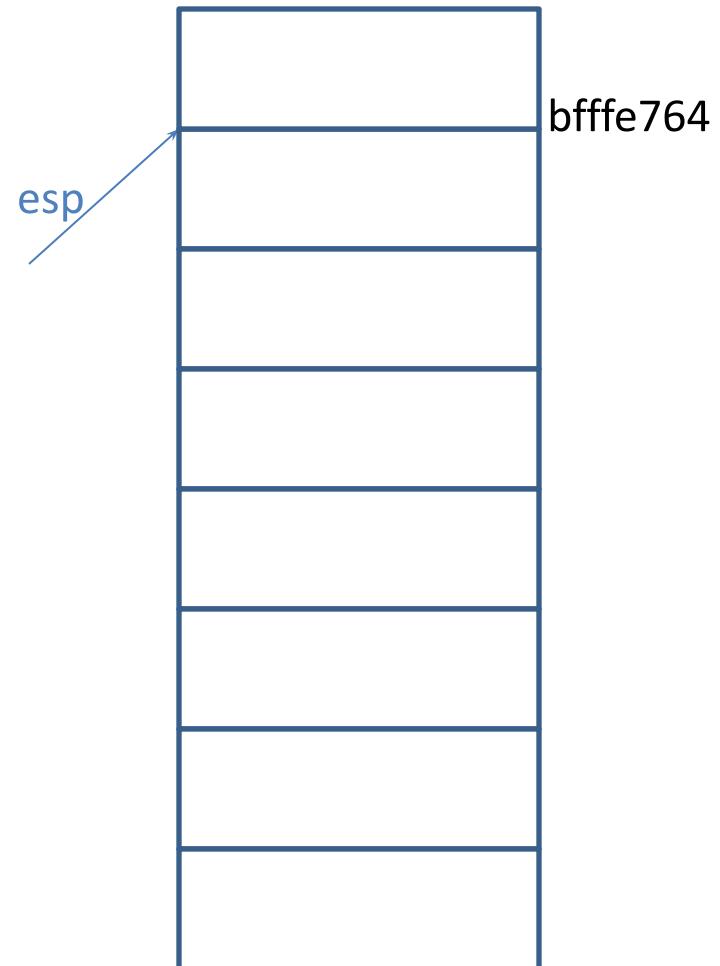
...

**main:**

...

```
subl $4, %esp  
movl $1, (%esp)  
call foo
```

...



...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp
```

...

**leave**

**ret**

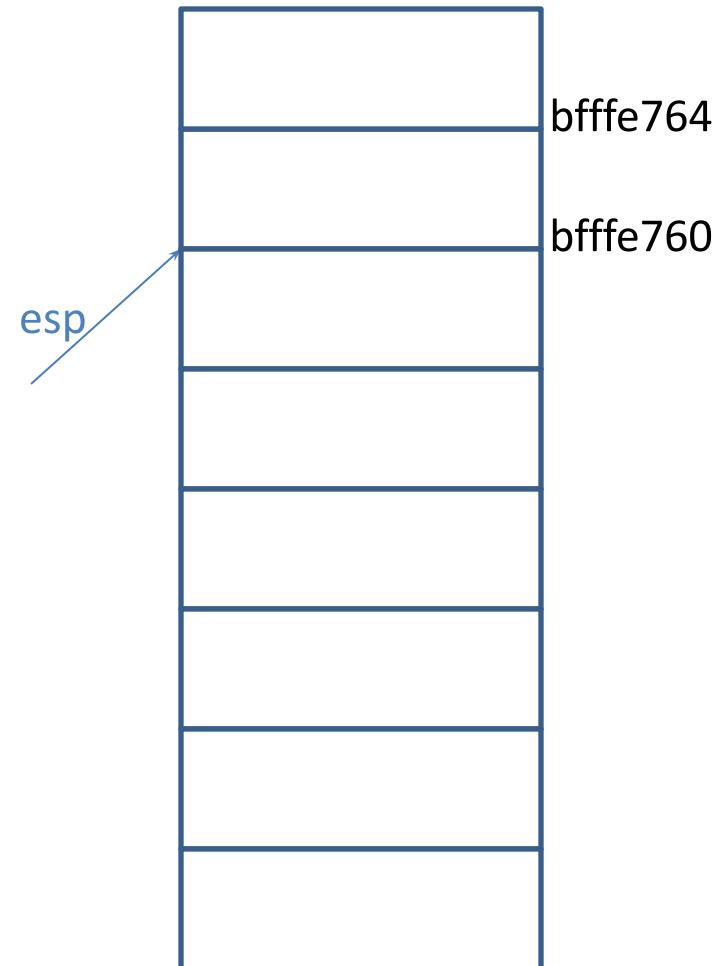
...

**main:**

...

```
subl $4, %esp  
movl $1, (%esp)  
call foo
```

...



...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp
```

...

**leave**

**ret**

...

**main:**

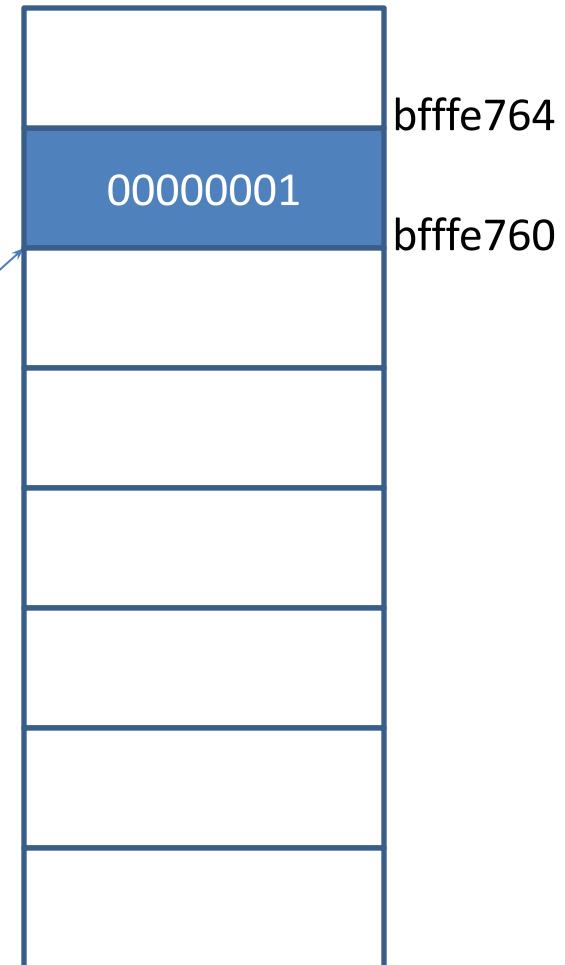
...

```
subl $4, %esp  
movl $1, (%esp)  
call foo
```

...

Arguments

esp



...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp
```

...

**leave**

**ret**

...

**main:**

...

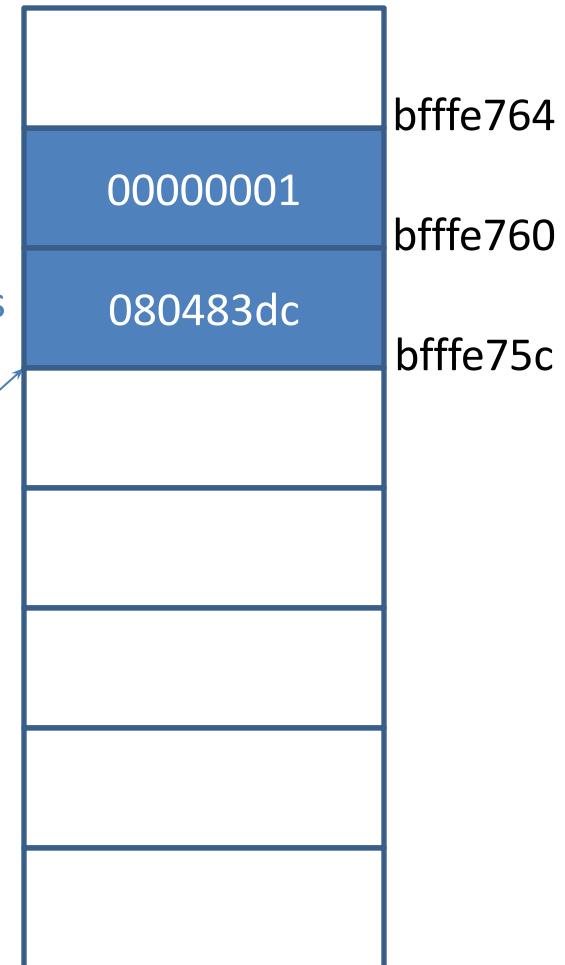
```
subl $4, %esp  
movl $1, (%esp)  
call foo
```

...

Arguments

Return address

esp



# Prologue

- A number of instructions are executed in the beginning of a function
- The idea is to prepare the stack and register for use within the function
- This process is given the name of *prologue*

# Prologue

- A number of instructions are executed in the beginning of a function
- The idea is to prepare the stack and register for use within the function
- This process is given the name of *prologue*
- Q: What happens during prologue?

# Prologue steps

1. Backs the old EBP up (e.g. function *main* EBP)
  - We push the old EBP onto the stack so we can restore its value later on after the function returns

# Prologue steps

1. Backs the old EBP up (e.g. function *main* EBP)
  - We push the old EBP onto the stack so we can restore its value later on after the function returns
2. Create the new EBP
  - We do this by making EBP point to the memory location of the current ESP

# Prologue steps

1. Backs the old EBP up (e.g. function *main* EBP)
  - We push the old EBP onto the stack so we can restore its value later on after the function returns
2. Create the new EBP
  - We do this by making EBP point to the memory location of the current ESP
3. Allocate the space needed for the function's local variables
  - By e.g. decrementing variables' sizes from ESP

...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp
```

...

**leave**

**ret**

...

**main:**

...

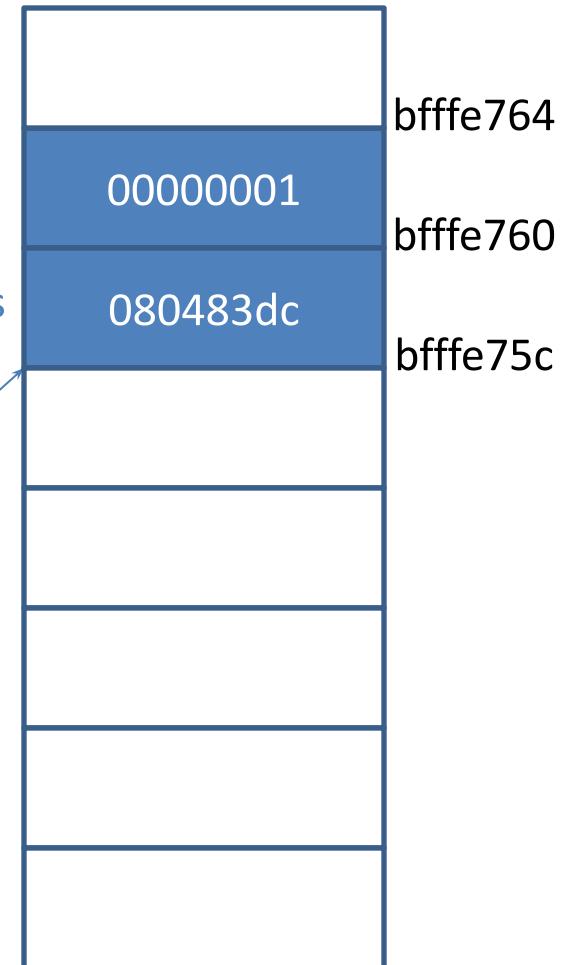
```
subl $4, %esp  
movl $1, (%esp)  
call foo
```

...

Arguments

Return address

esp



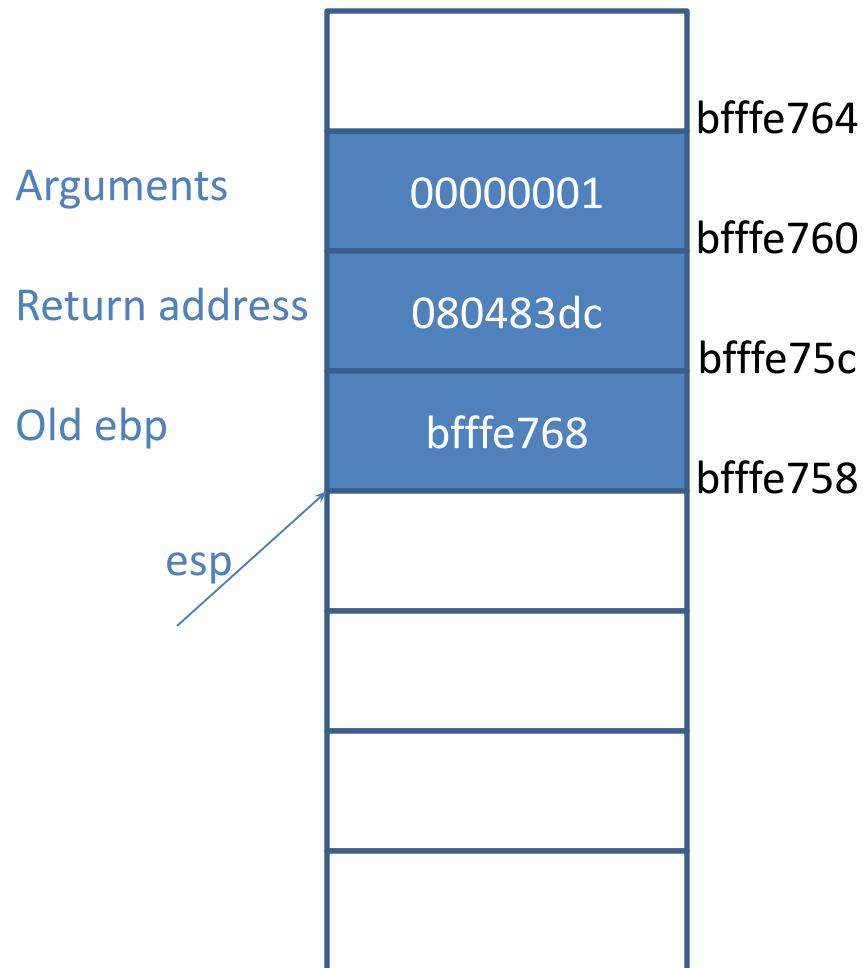
...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp  
...  
leave  
ret  
...
```

**main:**

```
...  
subl $4, %esp  
movl $1, (%esp)  
call foo  
...
```



...

**foo:**

**pushl %ebp**  
**movl %esp, %ebp**

**subl \$8, %esp**

...

**leave**

**ret**

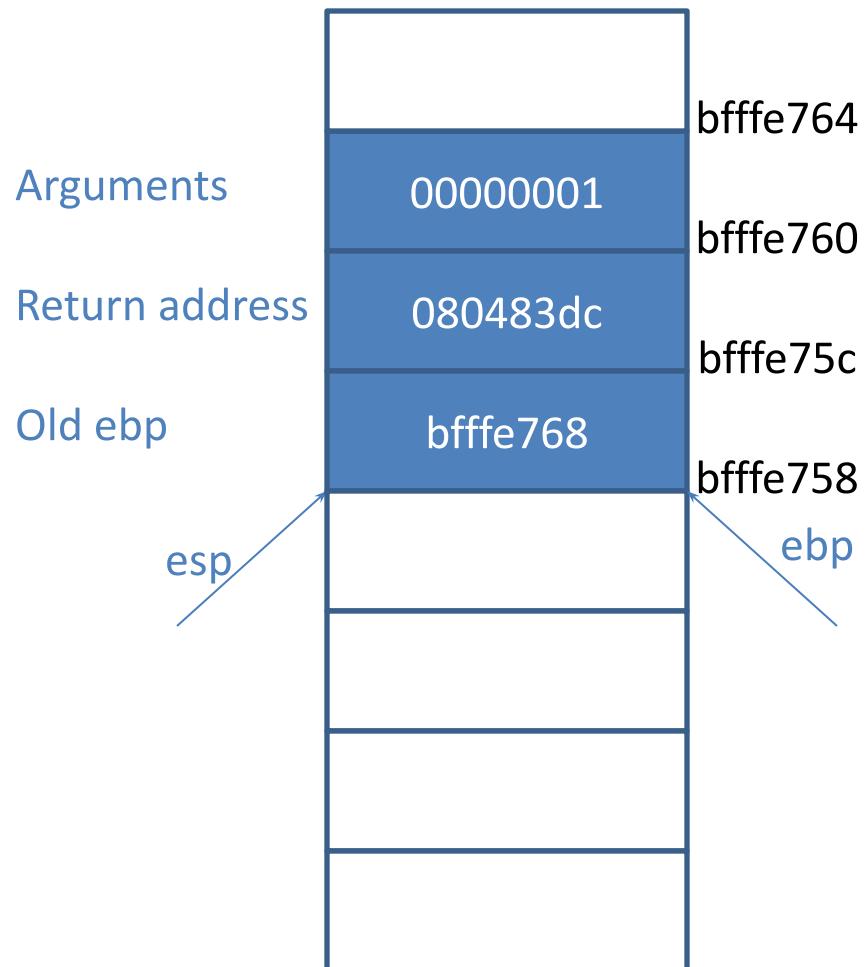
...

**main:**

...

**subl \$4, %esp**  
**movl \$1, (%esp)**  
**call foo**

...



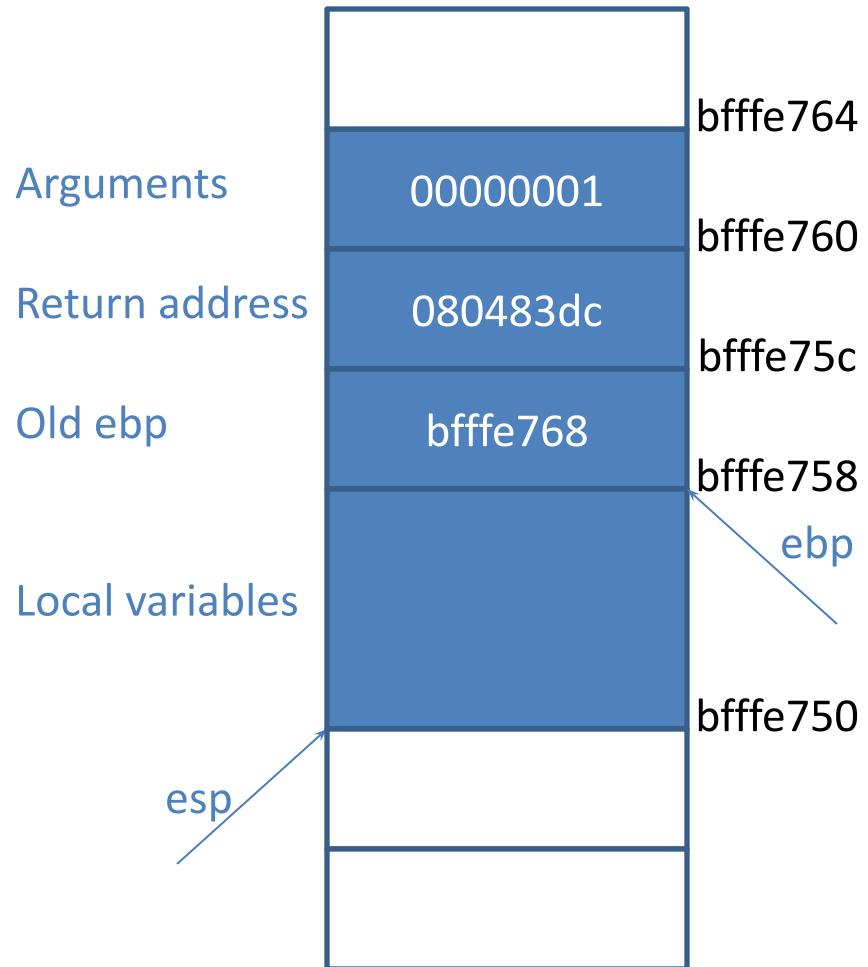
...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp  
...  
leave  
ret  
...
```

**main:**

```
...  
subl $4, %esp  
movl $1, (%esp)  
call foo  
...
```



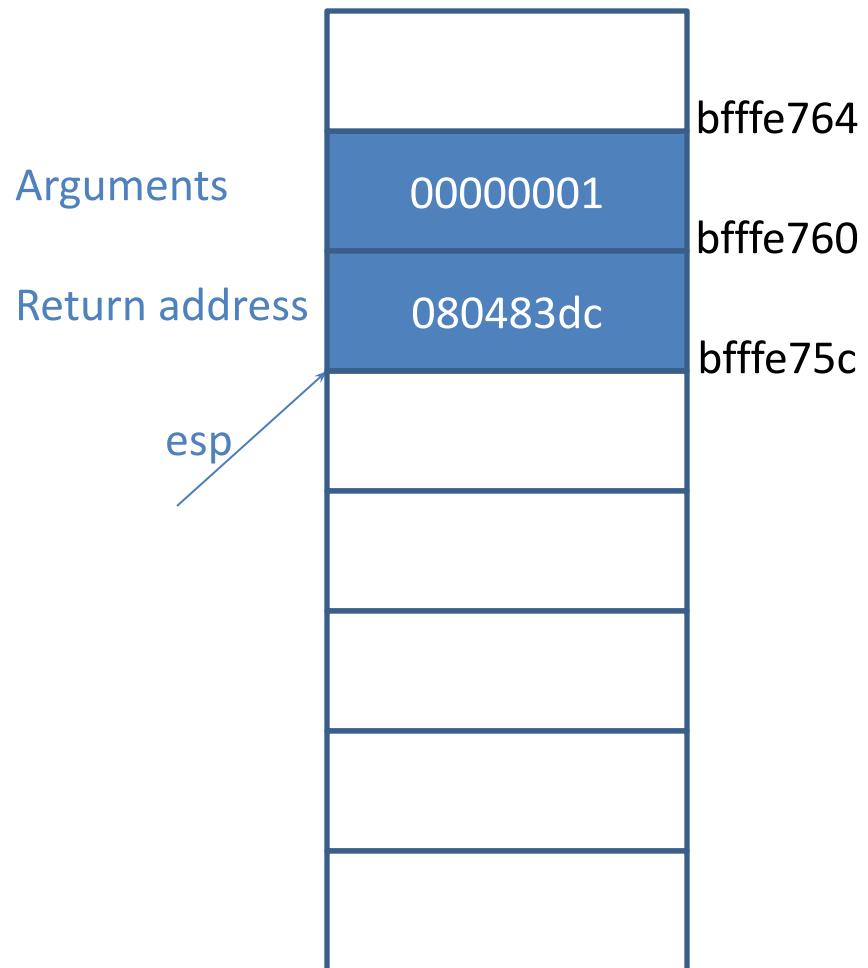
...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp  
...  
leave  
ret  
...
```

**main:**

```
...  
subl $4, %esp  
movl $1, (%esp)  
call foo  
...
```



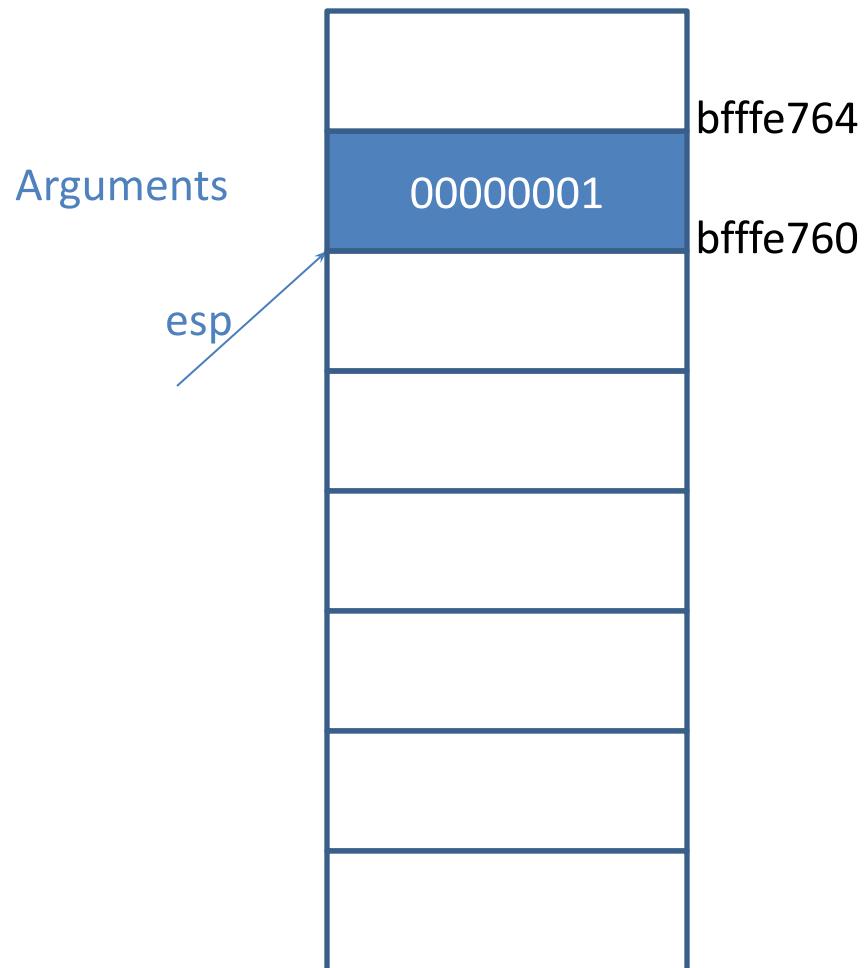
...

**foo:**

```
pushl %ebp  
movl %esp, %ebp  
subl $8, %esp  
...  
leave  
ret  
...
```

**main:**

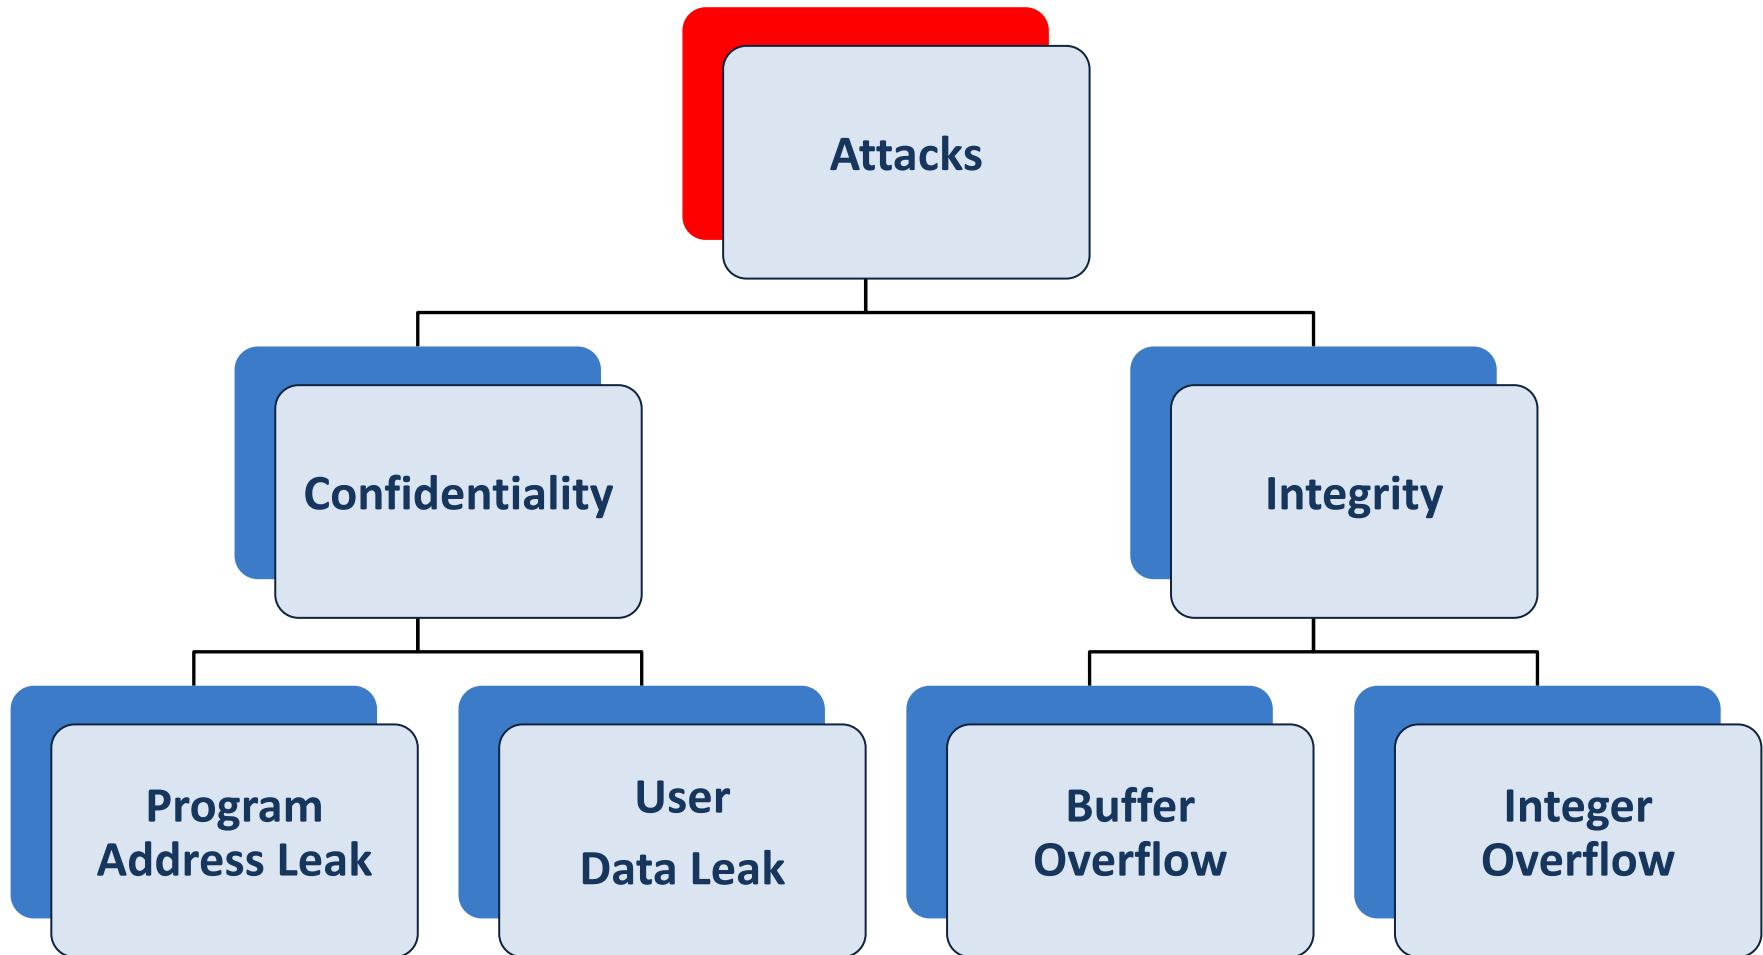
```
...  
subl $4, %esp  
movl $1, (%esp)  
call foo  
...
```



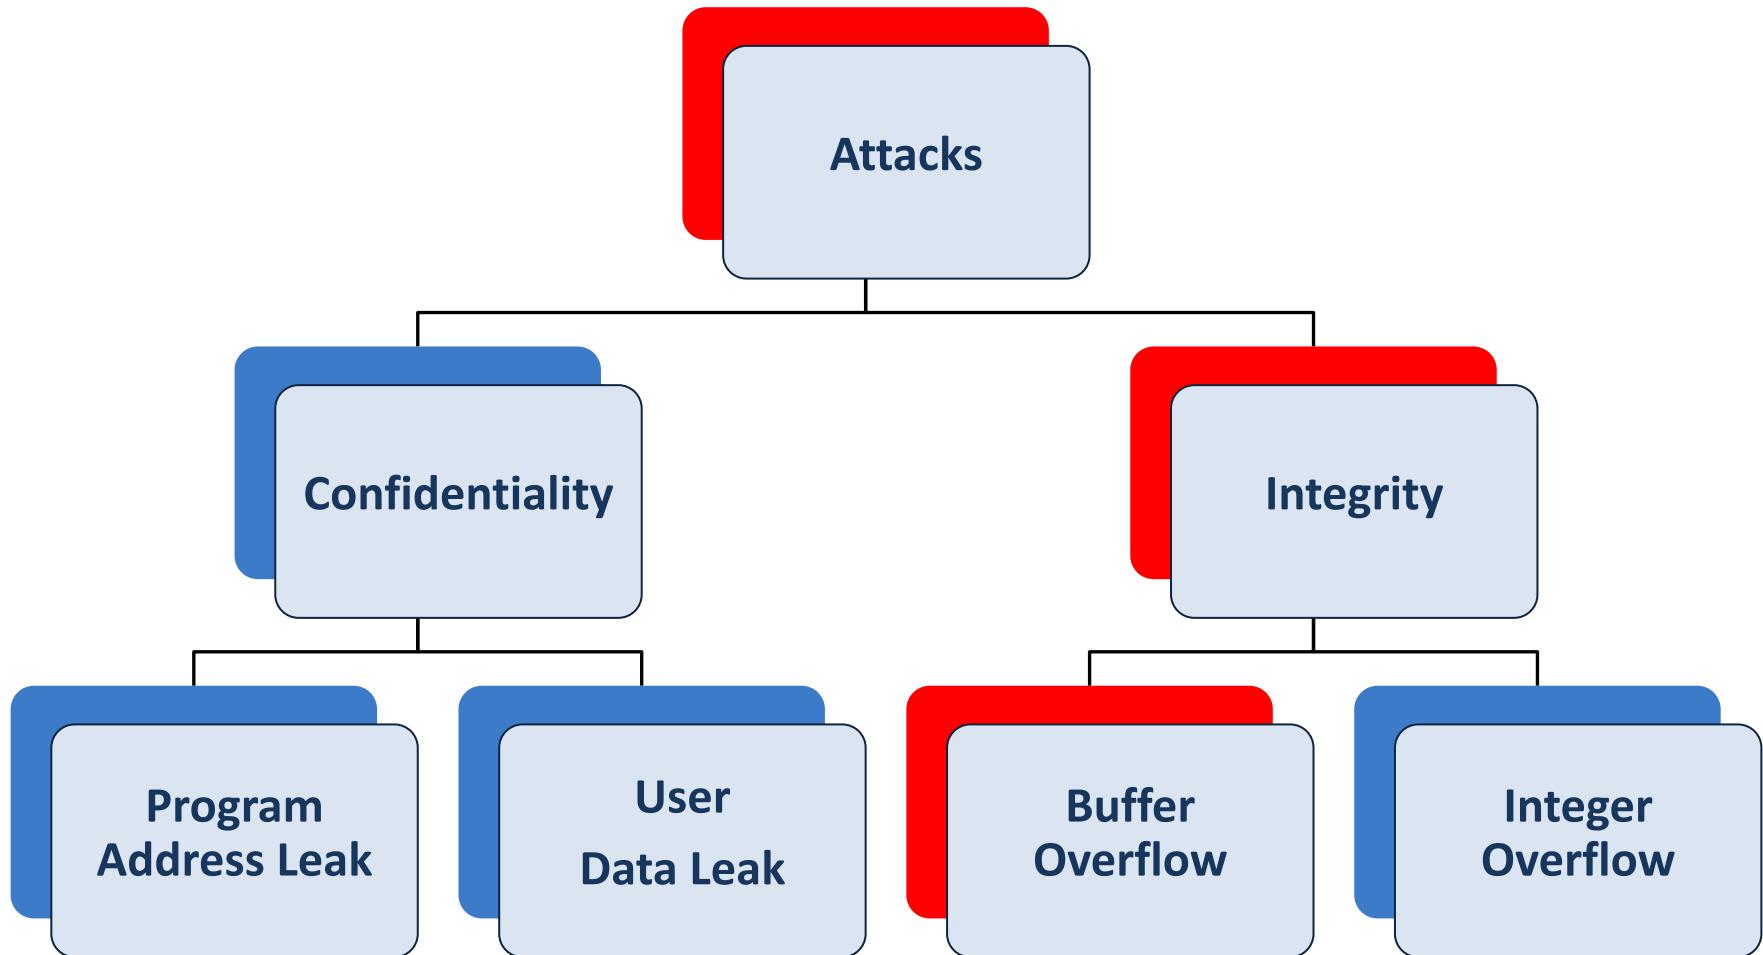
# Agenda

- C language
- Memory Management
- **Attacks**
- Countermeasures
- Final Remarks

# Attacks Taxonomy



# Attacks Taxonomy

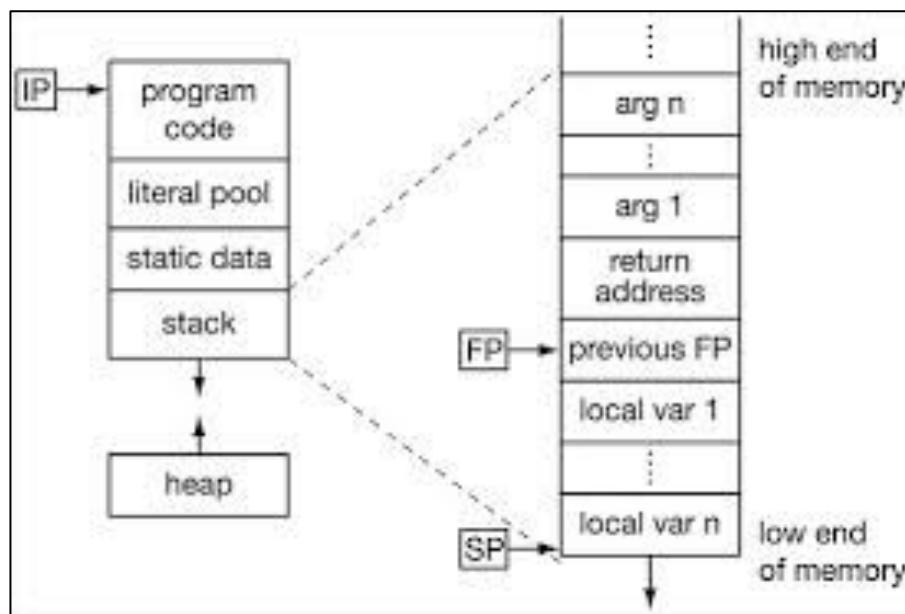


# Buffer Overflow

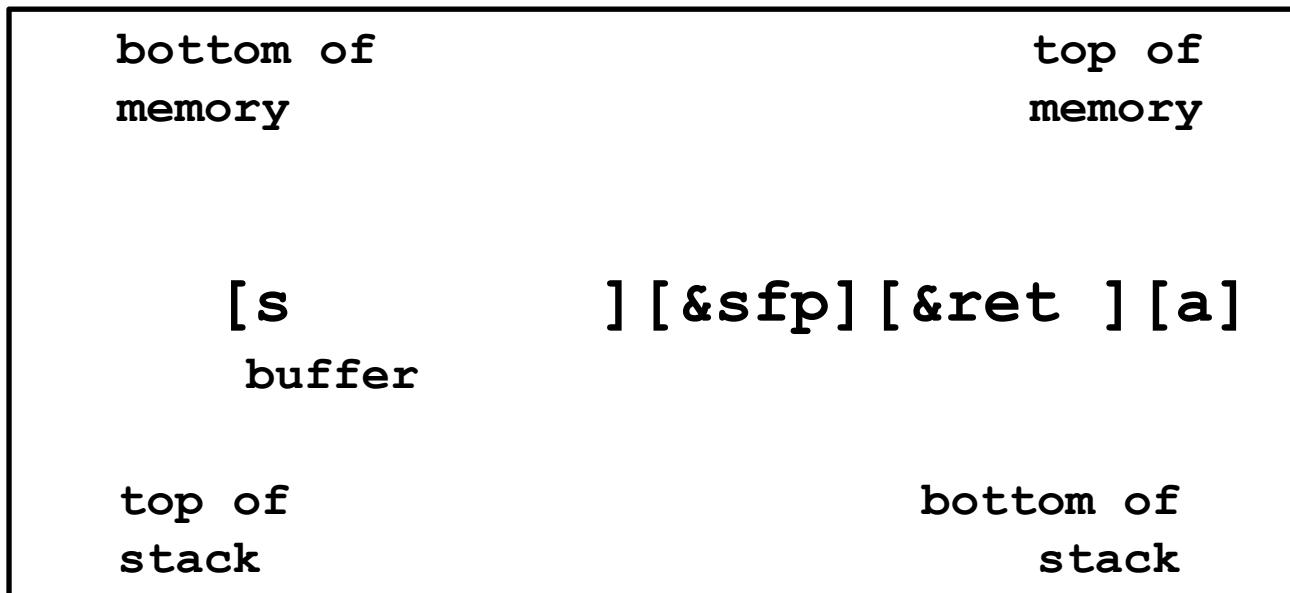
- A buffer overflow happens if, when writing to a buffer, access beyond its boundaries is made and adjacent memory addresses are modified
- The idea is to manipulate data through arrays that are not bound-checked
  - E.g. `strcpy(buffer, str);`

# Stack Smashing

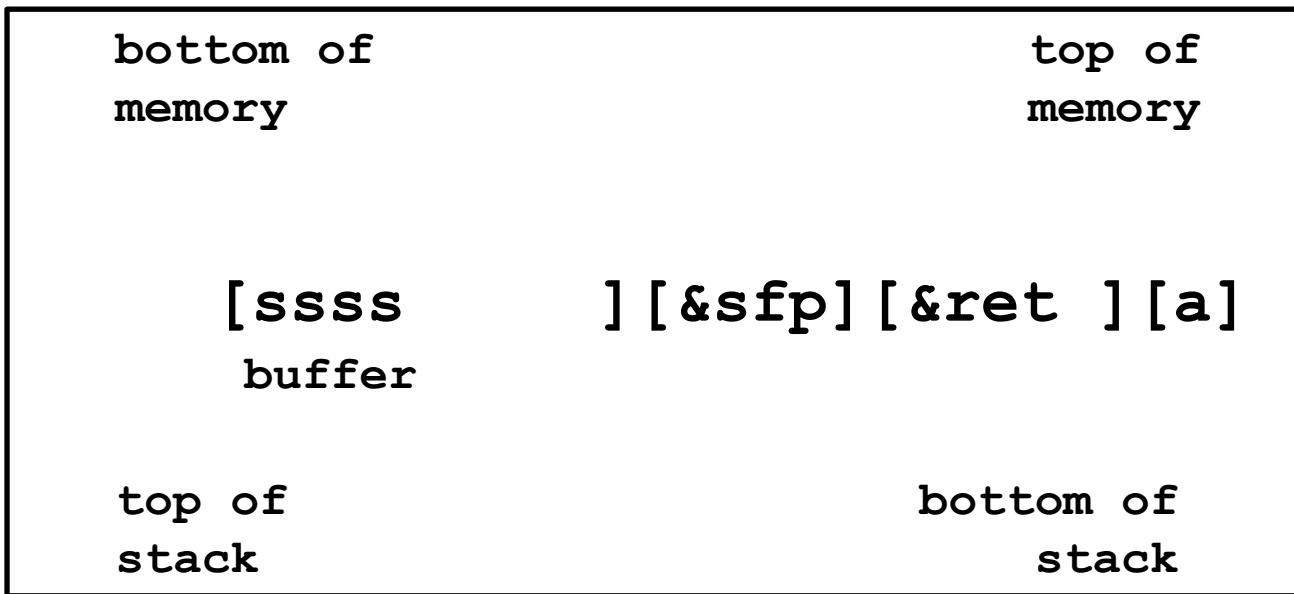
- *Stack smashing* is called when a buffer overflow happens inside a stack



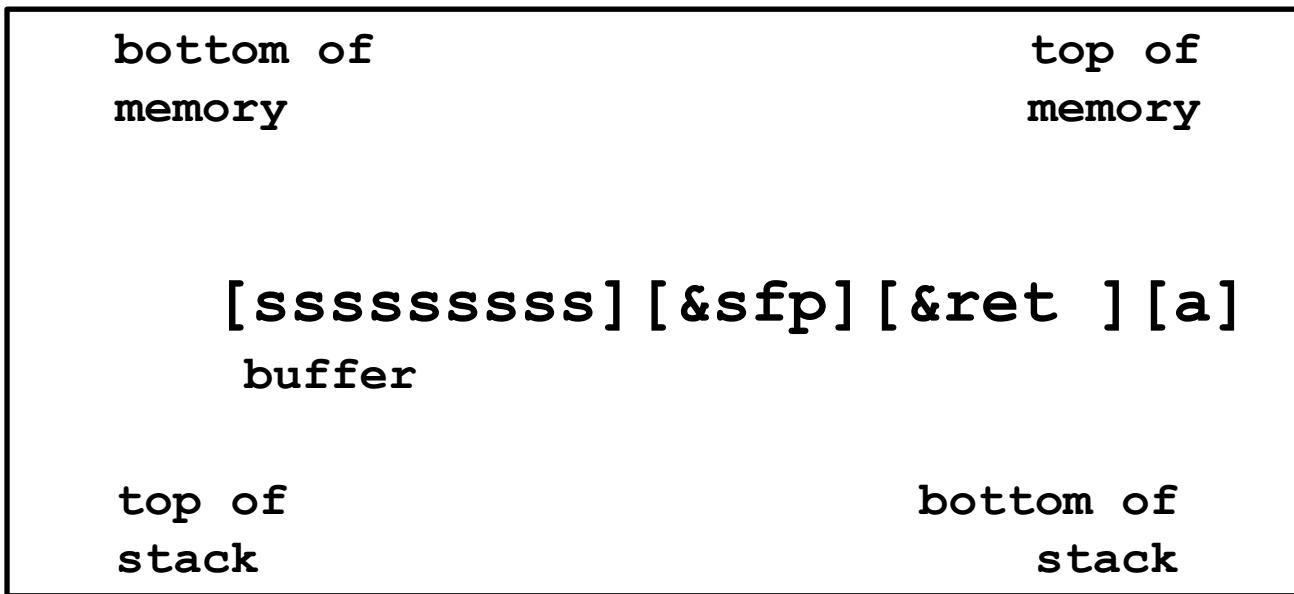
# Buffer Overflow Attack (Cont.)



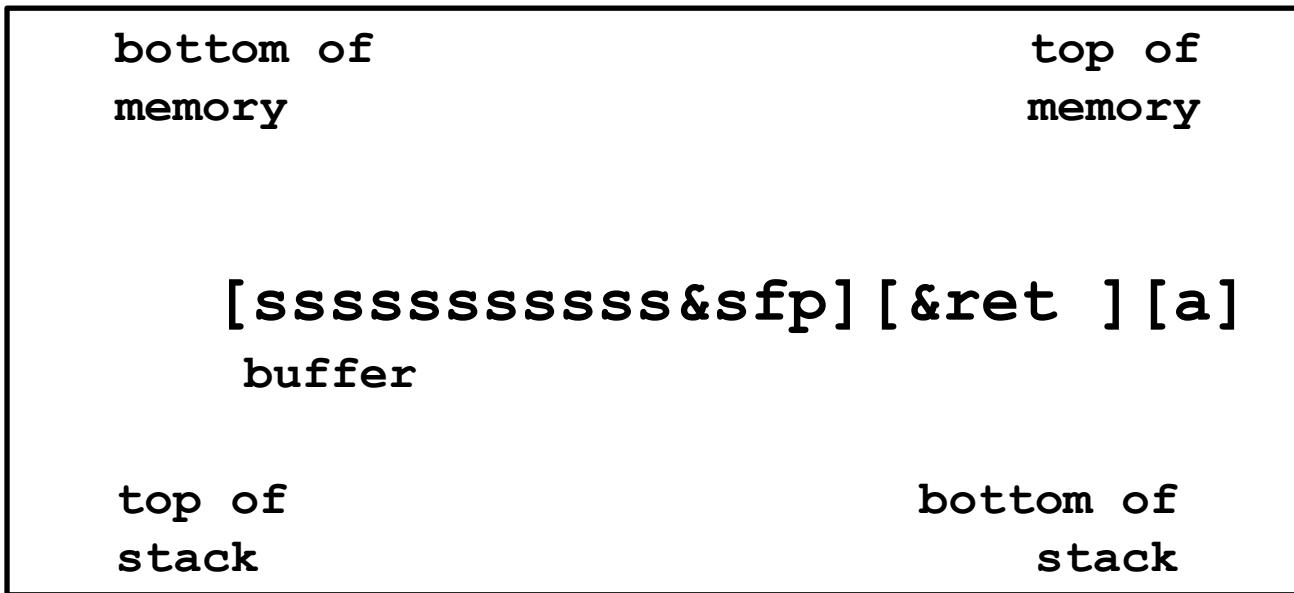
# Buffer Overflow Attack (Cont.)



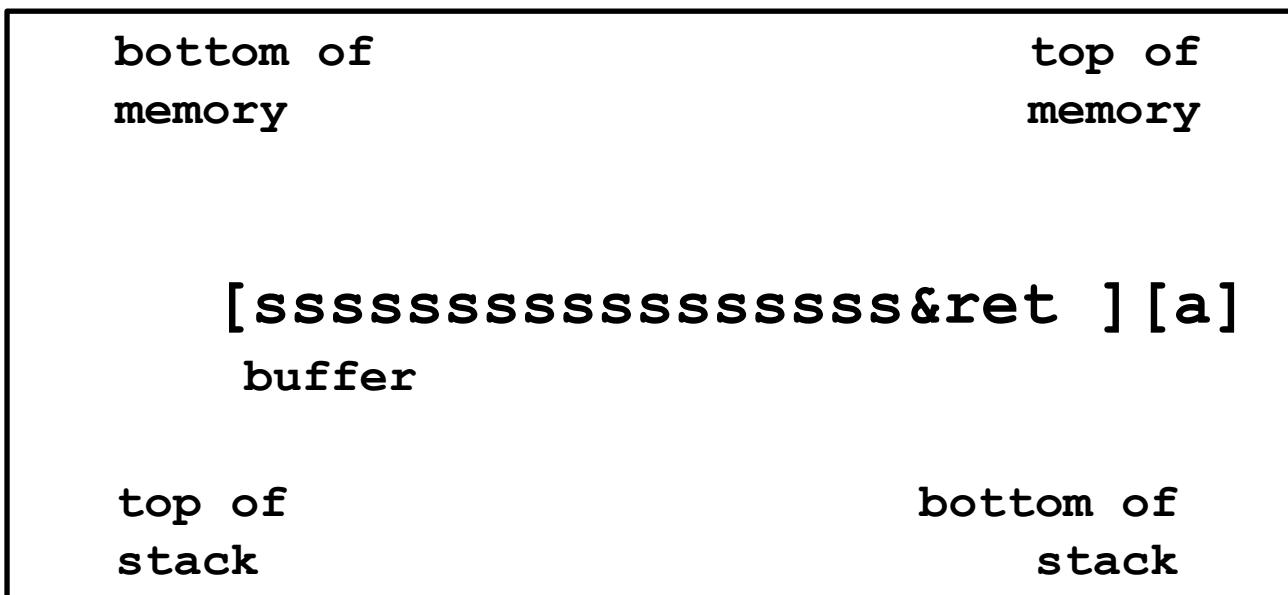
# Buffer Overflow Attack (Cont.)



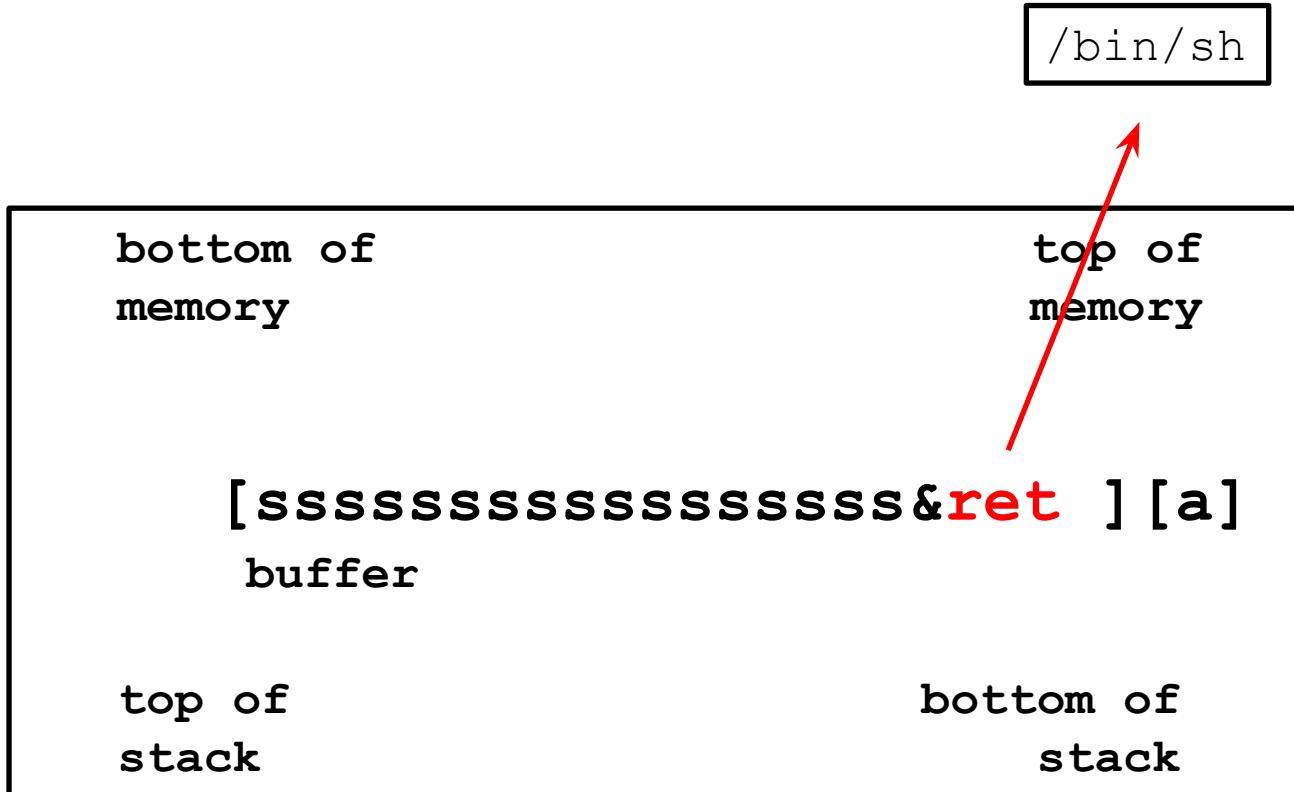
# Buffer Overflow Attack (Cont.)



# Buffer Overflow Attack (Cont.)



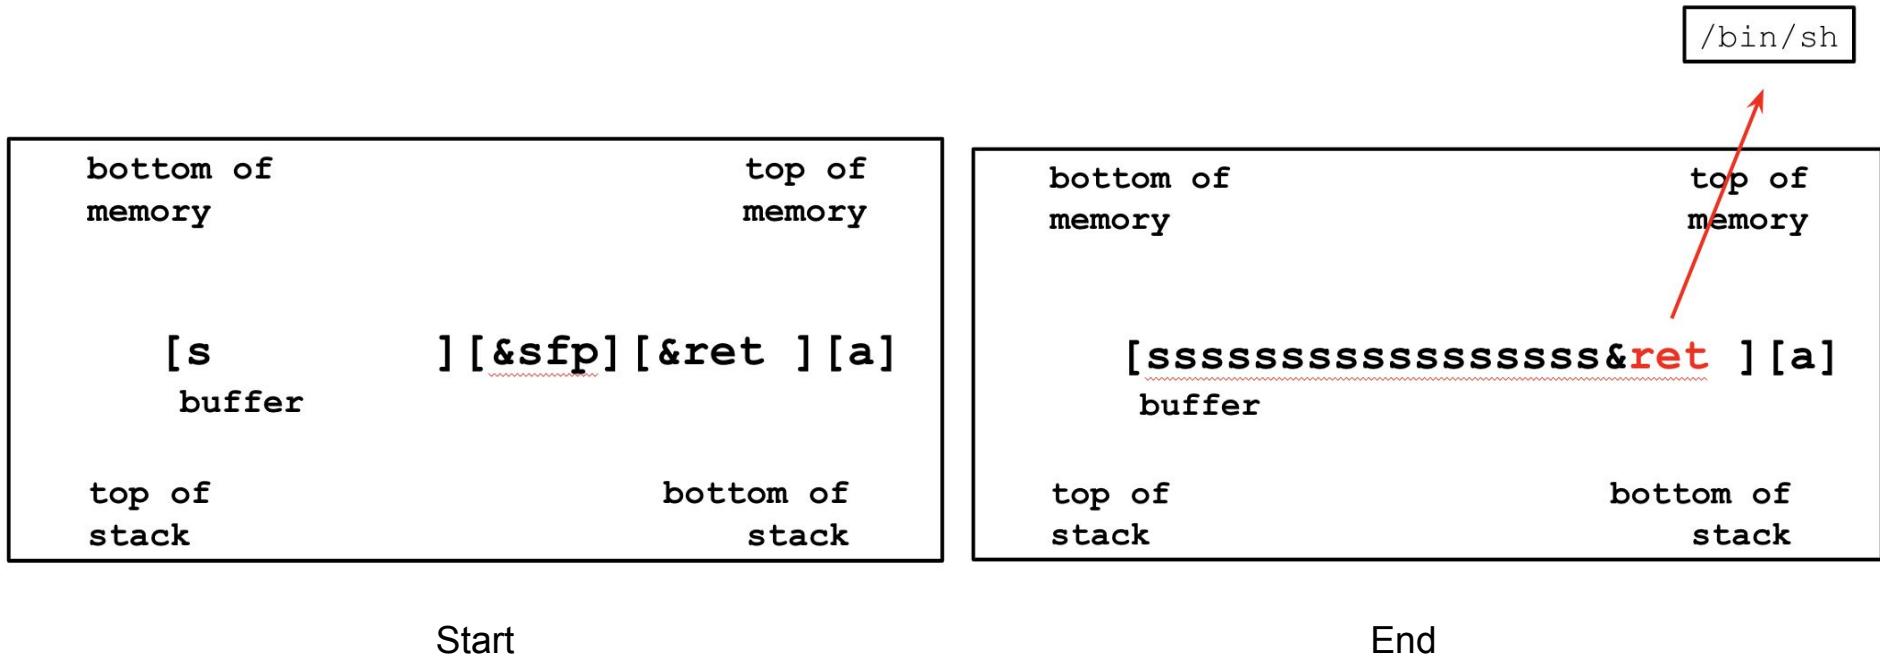
# Buffer Overflow Attack (Cont.)



# Agenda

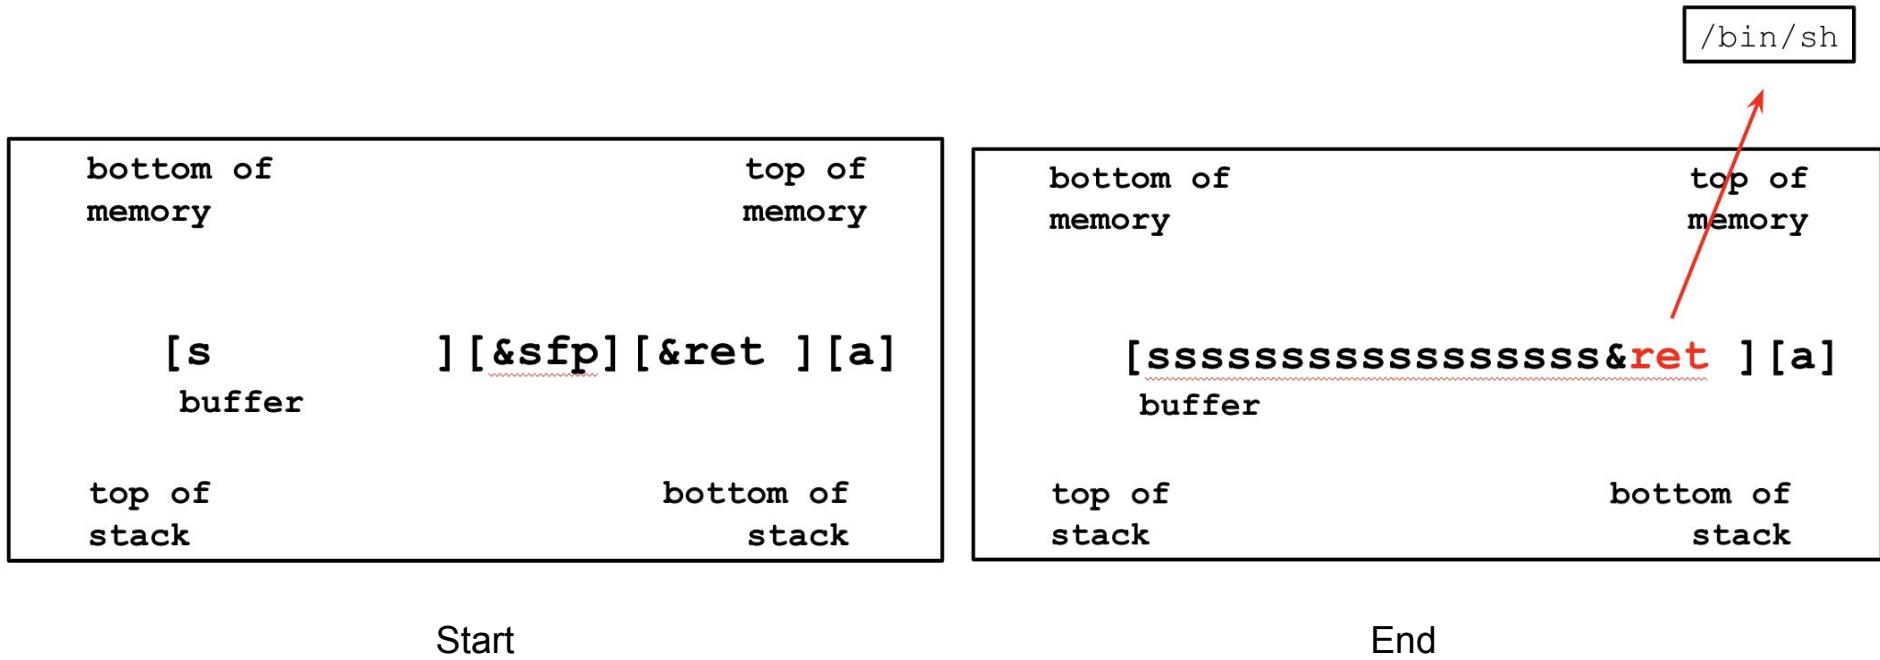
- C language
- Memory Management
- Attacks
- **Countermeasures**
- Final Remarks

# Buffer Overflow Attack



Q: Given the layouts before and after a BOF attack, do you have any ideas on how to solve this attack?

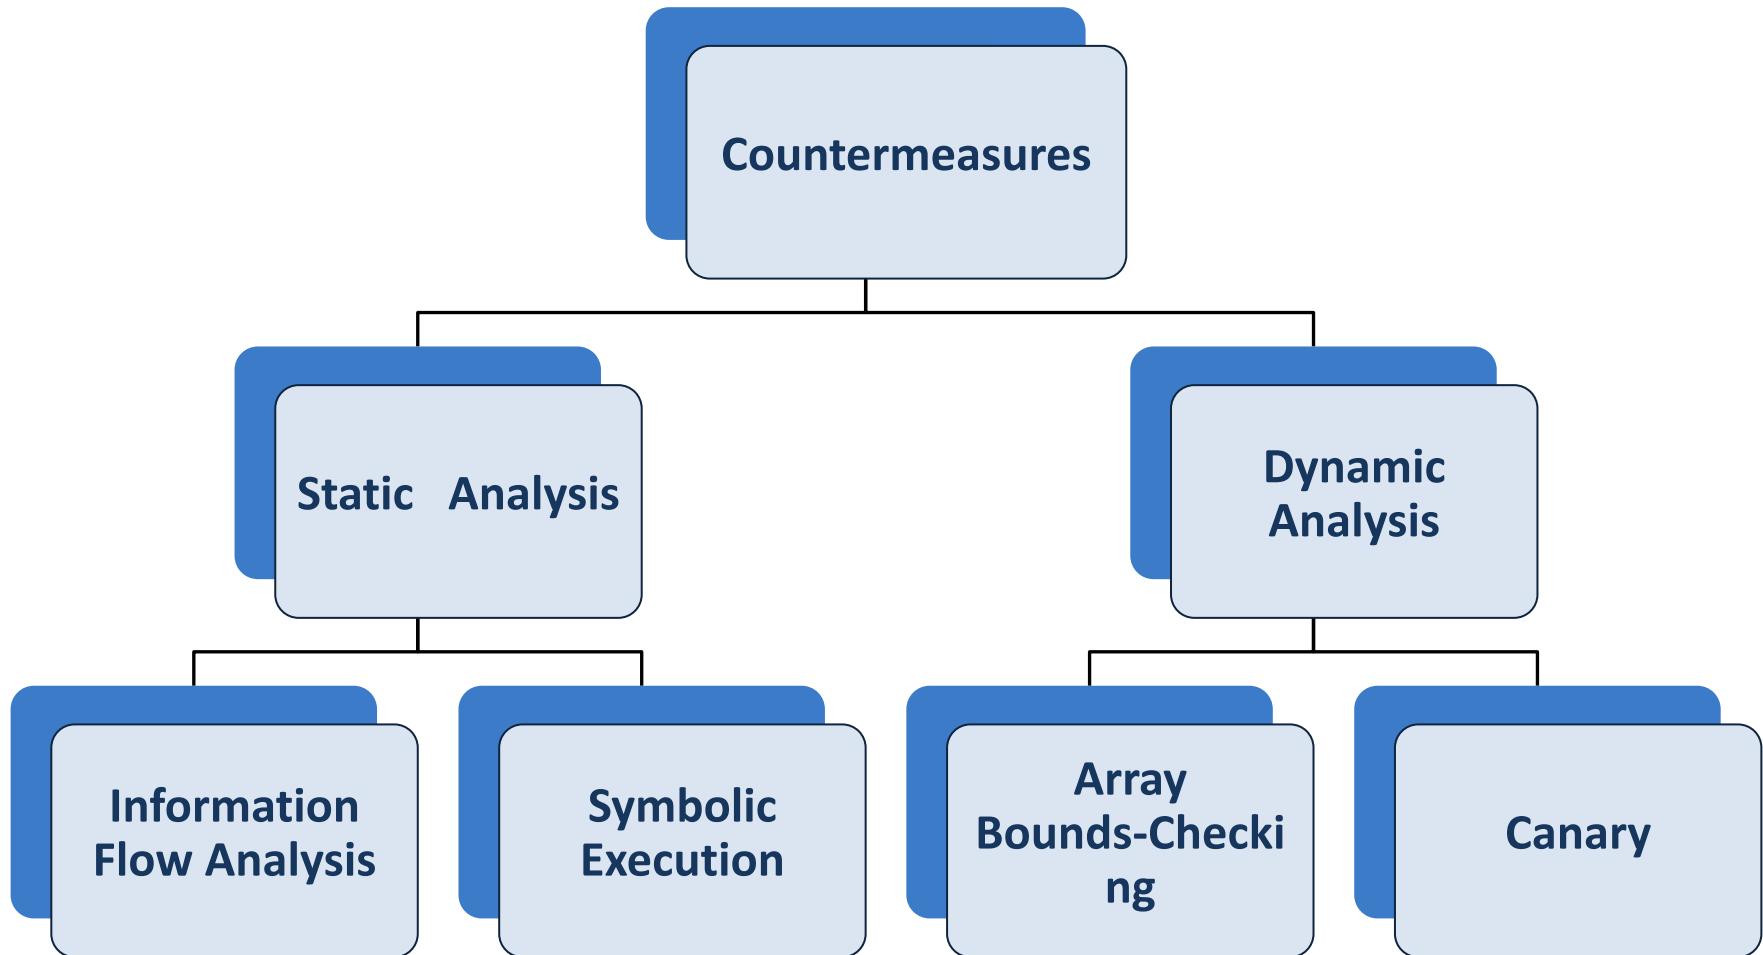
# Buffer Overflow Attack



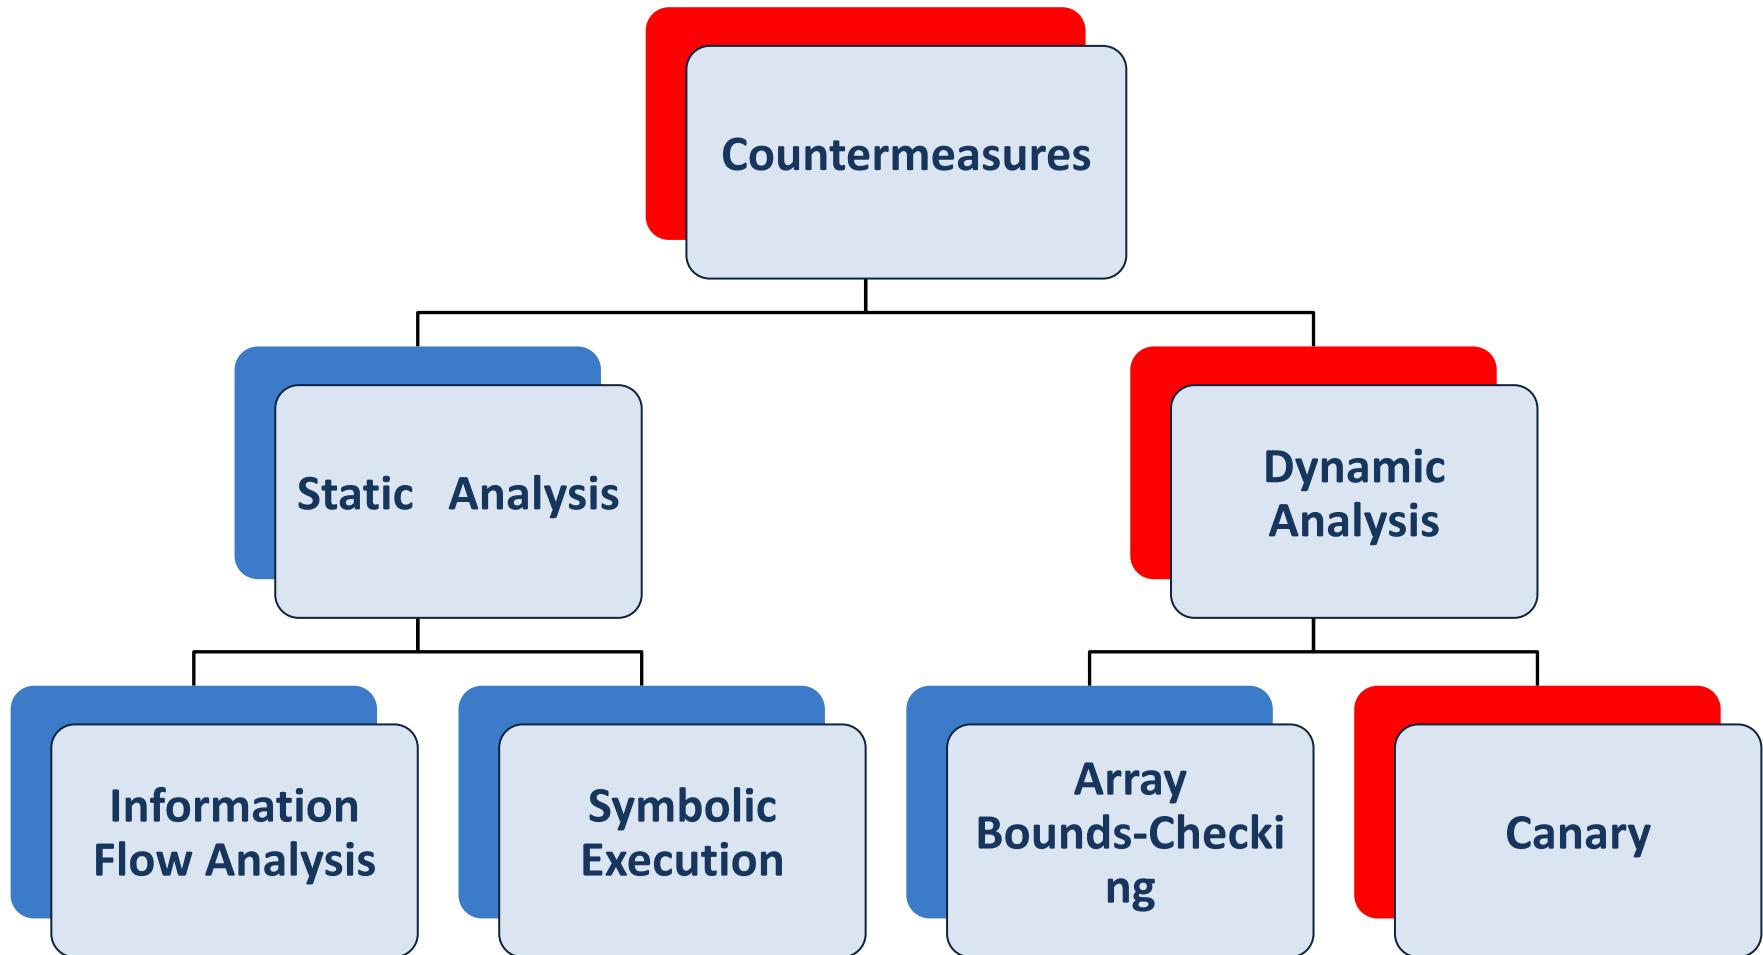
Q: Given the layouts before and after a BOF attack, do you have any ideas on how to solve this attack?

A: There are many, but we will look in to Canaries

# Countermeasures Taxonomy



# Countermeasures Taxonomy



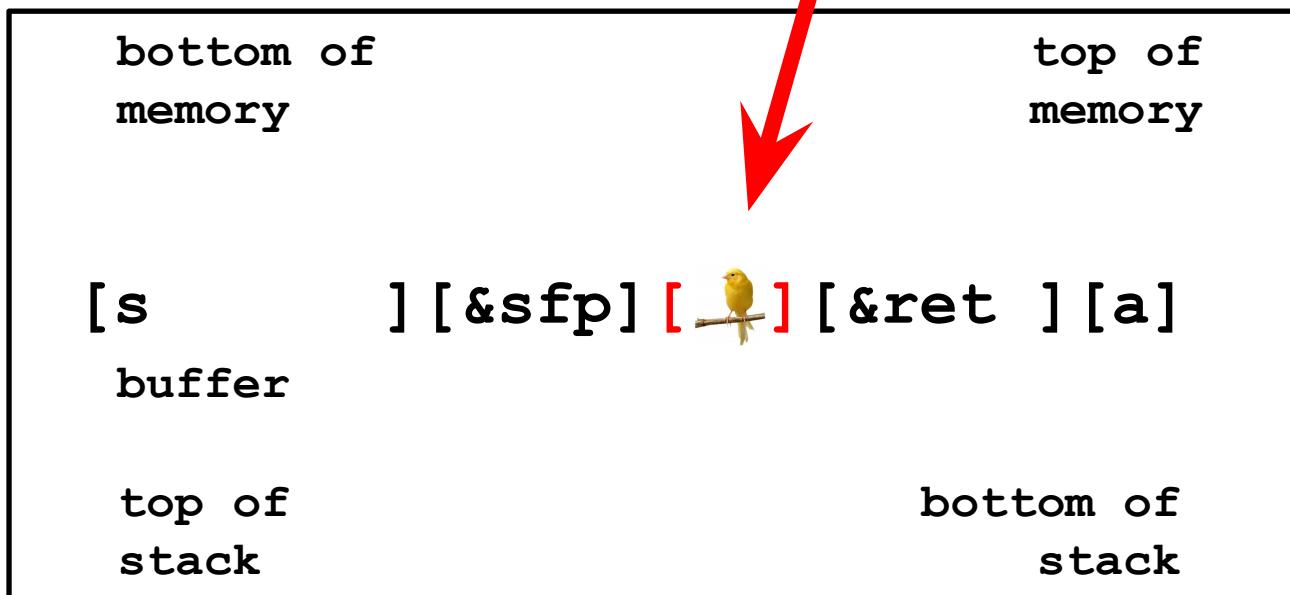


# Canary-based Protection

- Canary-based protection is one of the main protections against stack smashing
  - Proposed by Cowan et al. 1998
  - Like real canaries that detect dangerous gases in coal mines, canaries detects stack-smashing
- The idea is to place a *canary* just below the return address (RET)
  - Therefore, if RET is overwritten so is the canary
  - An alarm will then go off



# Canary-based Protection





# Canary: Pros & Cons

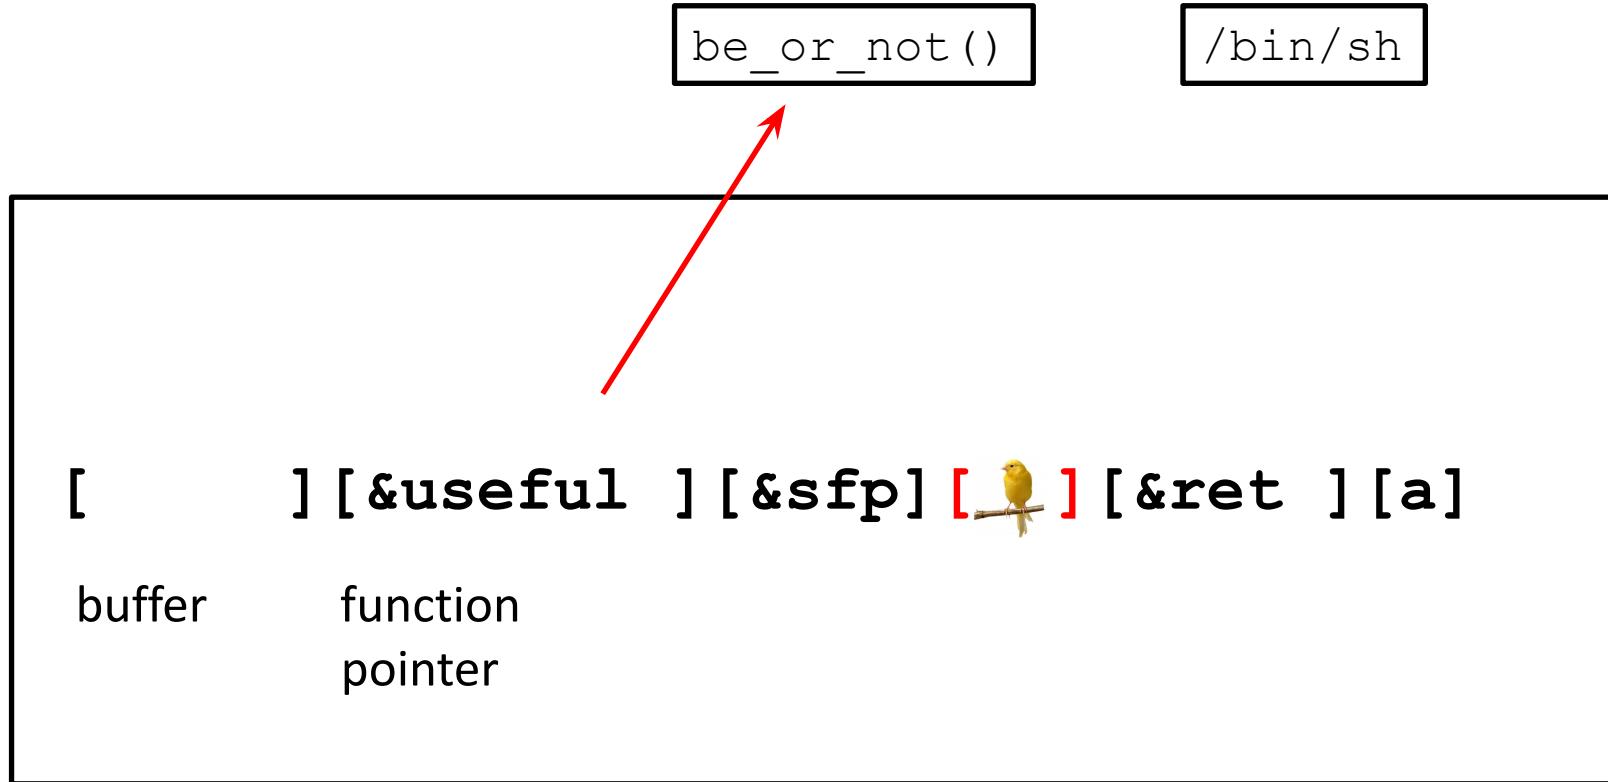
- Computationally cheap



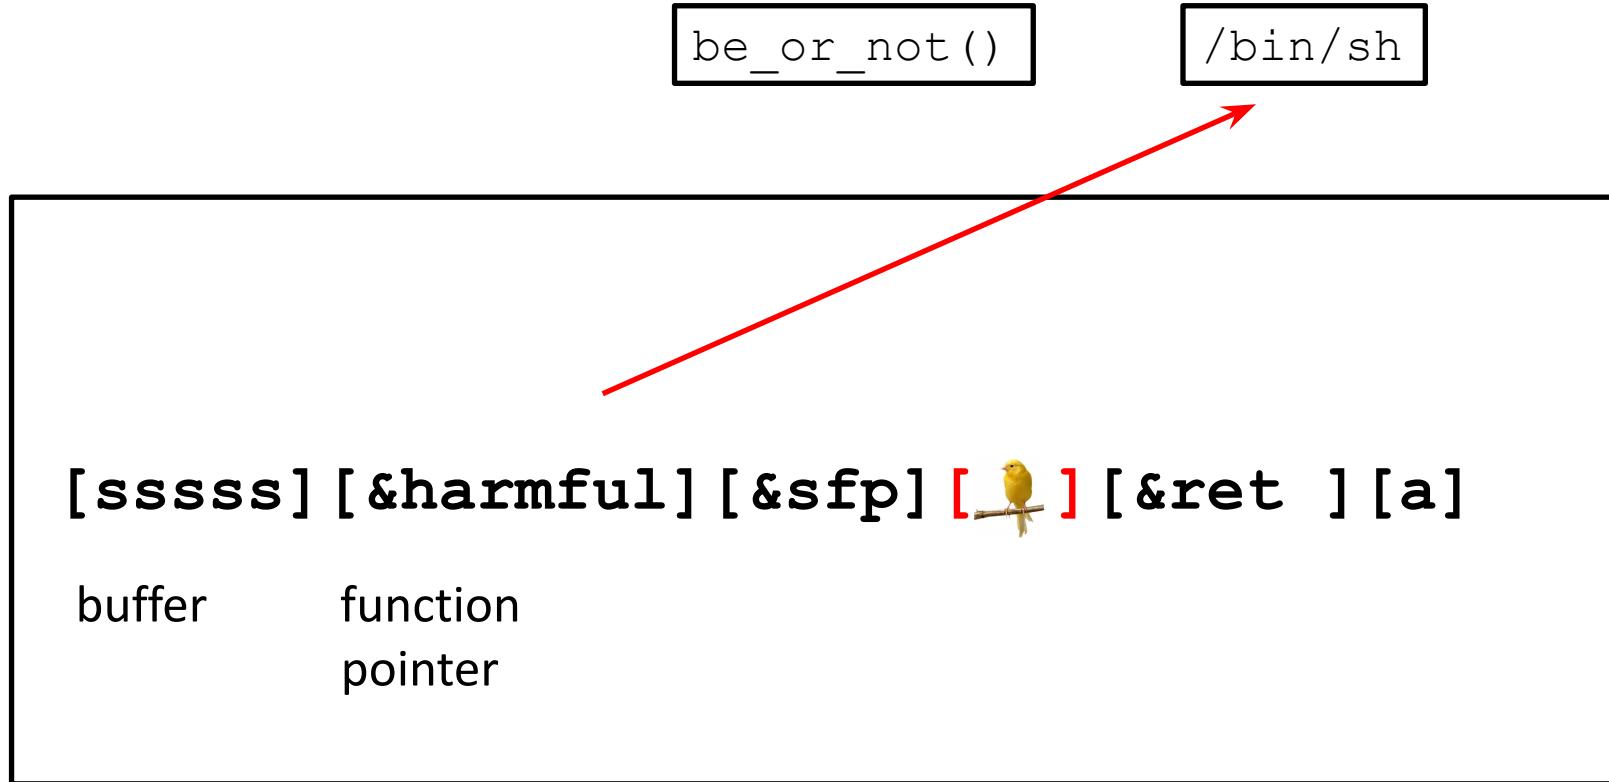
# Canary: Pros & Cons

- Computationally cheap
- Problem
  - Current canary-based protection proposals do not guard every kind of stack vulnerability

# Canary-based Protected Code Exploitation



# Canary-based Protected Code Exploitation



# Agenda

- C language
- Memory Management
- Attacks
- Countermeasures
- **Final Remarks**

# Final Remarks

- Security is all about getting right the dynamics of a computer
  - It has nothing to do with teenagers sitting in front of their screens and playing games
- Understating the security properties, attacks, and countermeasures will turn you into a security expert
- Security is wide field and the class today aimed to give you only an overview

# Final Remarks

- To convince you further about the relevance of this field, I want you students to watch this video
  - Stuxnet

# Questions?